

# Exercise 1

TMA4300 Computer Intensive Statistical Models

Mads Adrian Simonsen, William Scott Grundeland Olsen

01 februar, 2021

### Problem A: Stochastic simulation by the probability integral transform and bivariate techniques

1.

2.

(a)

(b)

3.

(a)

(b)

(c)

4.

5

### Problem B: The gamma distribution

1.

(a)

(b)

2.

(a)

(b)

3.

(a)

(b)

4.

5.

(a)

(b)

### Problem C: Monte Carlo integration and variance reduction

1.

2.

3.

(a)

(b)

### Problem D: Rejection sampling and importance sampling

#### Subproblem 1.

2

We consider a vector of multinomially distributed counts

$$\begin{bmatrix} y_1 \\ \vdots \end{bmatrix}$$

$$\begin{bmatrix} \frac{1}{2} + \frac{\theta}{4} \\ 1 - \theta^{\frac{1}{4}} \end{bmatrix}$$

and the observed data is  $\mathbf{y} = [125 \ 18 \ 20 \ 34]^\top$ . The multinomial mass function is given as

$$f(\mathbf{y} \mid \theta) \propto (2 + \theta)^{y_1} (1 - \theta)^{y_2 + y_3} \theta^{y_3},$$

and assuming a prior that is  $\text{Uniform}(0, 1)$  the posterior will be

$$f(\theta \mid \mathbf{y}) \propto f^*(\theta) := (2 + \theta)^{y_1} (1 - \theta)^{y_2 + y_3} \theta^{y_3},$$

for  $\theta \in (0, 1)$ . We wish to sample from this using a  $\text{Uniform}(0, 1)$  proposal density, that is,  $g(\theta \mid \mathbf{y}) = 1$ , for  $\theta \in (0, 1)$ . Because we do not know the normalizing constant, we may use weighted resampling, which is an approximate algorithm. We do this by first generating  $\Theta_1, \dots, \Theta_n \sim g(\theta) \sim \text{Uniform}(0, 1)$  and then calculate the weights

$$w(\Theta_i) = \frac{f(\Theta_i)/g(\Theta_i)}{\sum_{j=1}^n f(\Theta_j)/g(\Theta_j)} = \frac{f(\Theta_i)}{\sum_{j=1}^n f(\Theta_j)},$$

because  $g(\theta) = 1$  for all  $\theta \in (0, 1)$ . We then generate a second sample of size  $m$  from the discrete distribution  $\{\Theta_1, \dots, \Theta_n\}$  with probabilities  $w(\Theta_1), \dots, w(\Theta_n)$ . This is done in the code block below, where we choose  $m$  such that  $n/m = 20$ .

```
posterior_f_star <- function(theta, y) {
  return((2 + theta)^(y[1]) * (1 - theta)^(y[2] + y[3]) * theta^(y[4]))
}

weighted_resampling_f <- function(n, y) {
  Theta <- runif(n) # Generate n Uniform(0, 1) variables
  f_star <- posterior_f_star(Theta, y) # Calculating the vector f_star
  W <- f_star / sum(f_star) # Calculating the weights
  m <- n / 20 # Using m such that n / m = 20
  X <- sample(Theta, size = m, prob = W) # Sample m values from Theta with probability W
  return(X)
}
```

## Subproblem 2.

Drawing  $\Theta_1, \dots, \Theta_M \sim f(\theta \mid \mathbf{y})$ , the Monte Carlo estimate of  $\mu = E(\theta \mid \mathbf{y})$  is

$$\hat{\mu} = \frac{1}{M} \sum_{i=1}^M \Theta_i.$$

We do this for  $M = 10000$  in the code block below. Figure 1 shows the result of this. We see the estimation of the posterior mean  $E(\theta \mid \mathbf{y})$  using Monte Carlo integration and numerical integration together with the theoretical posterior density distribution and a generated histogram of the samples. In the figure the posterior density is plotted using a normalizing constant we find by numerical integration in R below, giving the normalizing constant `norm_const`.

```
set.seed(69)

M <- 10000 # Number of samples from f(theta | y)
n <- 20 * M # Number of samples needed to give M samples from f(theta | y)
y <- c(125, 18, 20, 34) # Observed data
Theta_samp <- weighted_resampling_f(n, y) # M samples from f(theta | y)
mu_est <- mean(Theta_samp) # = 1/M * sum(Theta_samp)

# Integrating numerically
norm_const <- integrate(function(theta) (posterior_f_star(theta, y)),
  lower = 0,
```

```

                                upper = 1)$value # Integrate to find normalizing constant
posterior_f <- function(theta, y) { # Creating the true posterior f
  return(posterior_f_star(theta, y) / norm_const)
}
mu_num <- integrate(function(theta)(theta * posterior_f(theta, y)),
                    lower = 0,
                    upper = 1)$value # Value of mu by numerical integration

# Plot
ggplot() +
  geom_histogram(
    data = as.data.frame(Theta_samp),
    mapping = aes(x = Theta_samp, y = ..density..),
    binwidth = 0.01,
    boundary = 0
  ) +
  stat_function(
    fun = posterior_f,
    args = list(y = y),
    aes(col = "Posterior density")
  ) +
  geom_vline(
    aes(xintercept = c(mu_est, mu_num),
        col = c("Estimated posterior mean", "Numerical posterior mean"),
        linetype = c("dashed", "dotted"))
  ) +
  guides(linetype = FALSE) + # Remove linetype from label
  ggtitle("Estimation of the posterior mean") +
  xlab("theta") +
  ylab("Density") +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5)) +
  theme(legend.title = element_blank())

```

In the following code block we find the values of `mu_est` and `mu_num`.

```
mu_est
```

```
## [1] 0.6225143
```

```
mu_num
```

```
## [1] 0.6228061
```

From this it is clear that the estimated posterior mean is  $\hat{\mu} \approx 0.623$  using Monte Carlo integration, and  $\mu \approx 0.623$  using numerical integration with `integrate()`. Figure 1 also shows that these means corresponds well to the real posterior mean.

### Subproblem 3.

### Subproblem 4.

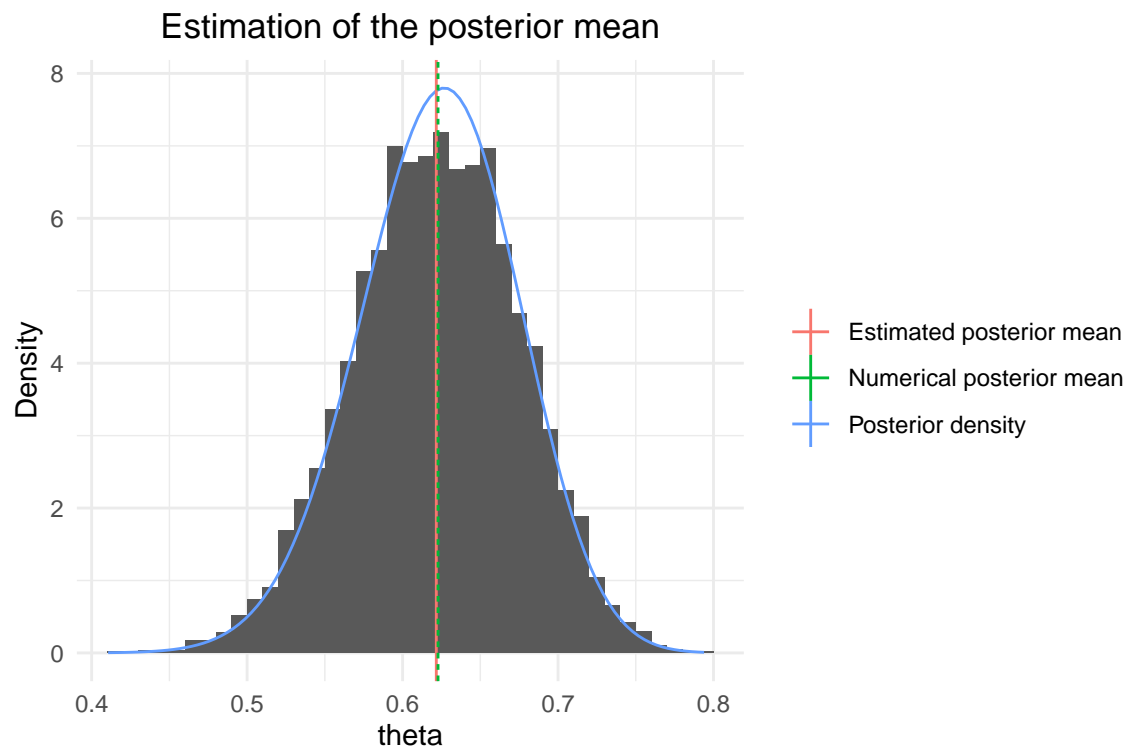


Figure 1: Estimation of the posterior mean  $E(\theta \mid \mathbf{y})$  using Monte Carlo integration and numerical integration. A histogram of the samples is also shown together with the theoretical posterior density distribution.