Exercise 1

TMA4300 Computer Intensive Statistical Models

Mads Adrian Simonsen, William Scott Grundeland Olsen

04 februar, 2021

Problem A: Stochastic simulation by the probability integral transform and bivariate techniques

Subproblem 1.

Let $X \sim \text{Exponential}(\lambda)$, with the cumulative density function

$$F_X(x) = 1 - e^{-\lambda x}, \quad x \geq 0.$$

Then the random variable $Y := F_X(X)$ has a $\text{Uniform}(0, 1)$ distribution. The probability integral transform becomes

$$Y = 1 - e^{-\lambda X} \quad \Leftrightarrow \quad X = -\frac{1}{\lambda} \ln(1 - Y).$$

It is clear that if $U \sim \text{Uniform}(0, 1)$, then $1 - U \sim \text{Uniform}(0, 1)$, and therefore we may as well say that

$$X = -\frac{1}{\lambda} \ln(Y). \tag{1}$$

Thus, we sample Y from `runif()` and transform it using Equation (1), to sample from the exponential distribution. Figure 1 shows one million samples drawn from the `generate_from_exp()` function defined in the code chunk below. It also shows the theoretical PDF of the exponential distribution with rate parameter $\lambda = 2$.

```
#set.seed(123)

generate_from_exp <- function(n, rate = 1) {
  Y <- runif(n)
  X <- -(1 / rate) * log(Y)
  return(X)
}

# sample
n <- 1000000   # One million samples
lambda <- 2
exp_samp <- generate_from_exp(n, rate = lambda)

# plot
ggplot() +
  geom_histogram(
```

```
    data = as.data.frame(exp_samp),
    mapping = aes(x = exp_samp, y = ..density..),
    binwidth = 0.05,
    boundary = 0
) +
stat_function(
    fun = dexp,
    args = list(rate = lambda),
    aes(col = "Theoretical density")
) +
ylim(0, lambda) +
xlim(0, 2) +
ggtitle("Simulating from an exponential distribution") +
xlab("x") +
ylab("Density") +
theme_minimal() +
theme(plot.title = element_text(hjust = 0.5))
```


Figure 1: Normalized histogram of one million samples drawn from the exponential distribution, together with the theoretical PDF, with $\lambda = 2$.

Theoretically, the mean and variance of $X \sim \text{Exponential}(\lambda)$ is $\text{E}(X) = \lambda^{-1}$ and $\text{Var}(X) = \lambda^{-2}$. So for $\lambda = 2$ we would expect $\text{E}(X) = 1/2$ and $\text{Var}(X) = 1/4$. For the simulation we get the mean and variance as calculated in the code block below, showing what we would expect.

```
mean(exp_samp)
```

```
## [1] 0.499487
```

```
var(exp_samp)
```

```
## [1] 0.24952
```

Subproblem 2.

Subsubproblem (a)

We are considering the probability density function

$$g(x) = \begin{cases} cx^{\alpha-1} & \text{if } 0 < x < 1, \\ ce^{-x} & \text{if } x \geq 1, \\ 0 & \text{otherwise}, \end{cases} \tag{2}$$

where c is a normalizing constant and $\alpha \in (0,1)$. If $x \leq 0$ the cumulative distribution function is zero. In the interval $0 < x < 1$ it becomes

$$G(x) = \int_{-\infty}^{x} g(\xi)\,\mathrm{d}\xi = \int_{0}^{x} c\xi^{\alpha-1}\,\mathrm{d}\xi = \frac{c}{\alpha}[\xi^{\alpha}]_{0}^{x} = \frac{c}{\alpha}x^{\alpha},$$

and finally for $x \geq 1$ we have

$$G(x) = \int_{-\infty}^{x} g(\xi)\,\mathrm{d}\xi = \int_{0}^{1} c\xi^{\alpha-1}\,\mathrm{d}\xi + \int_{1}^{x} ce^{-\xi}\,\mathrm{d}\xi = \left[\frac{c}{\alpha}\xi^{\alpha}\right]_{0}^{1} - \left[ce^{-\xi}\right]_{1}^{x} = c\left(\frac{1}{\alpha} - e^{-x} + \frac{1}{e}\right),$$

for $\alpha \in (0,1)$. That is, the cumulative density function is

$$G(x) = \begin{cases} 0 & \text{if } x \leq 0, \\ \frac{c}{\alpha}x^{\alpha} & \text{if } 0 < x < 1, \\ c\left(\frac{1}{\alpha} - e^{-x} + \frac{1}{e}\right) & \text{if } x \geq 1. \end{cases}$$

In this case it is trivial to find c. We solve

$$1 = \int_{\mathbb{R}} g(x)\,\mathrm{d}x = \int_{0}^{1} cx^{\alpha-1}\,\mathrm{d}x + \int_{1}^{\infty} ce^{-x}\,\mathrm{d}x = \frac{c}{\alpha} + \frac{c}{e},$$

which gives that

$$c = \frac{\alpha e}{\alpha + e}.$$

Writing the cumulative density function using this as c we obtain

$$G(x) = \begin{cases} 0 & \text{if } x \leq 0, \\ \frac{e}{\alpha+e}x^{\alpha} & \text{if } 0 < x < 1, \\ 1 - \frac{\alpha}{\alpha+e}e^{1-x} & \text{if } x \geq 1, \end{cases}$$

for $\alpha \in (0,1)$.

We may then find the inverse cumulative function. For $x \leq 0$ this is just zero, and for $0 < x < 1$, that is $0 < G(x) < \frac{e}{\alpha+e}$, we solve $x = \frac{e}{\alpha+e}y^{\alpha}$ for y giving $G^{-1}(x) = \left(\frac{\alpha+e}{e}x\right)^{1/\alpha}$. Similarly for $x \geq 1$, that is $G(x) \geq 1 - \frac{\alpha}{\alpha+e} = \frac{e}{\alpha+e}$, we solve $x = 1 - \frac{\alpha}{\alpha+e}e^{1-y}$ for y, such that

$$G^{-1}(x) = \begin{cases} \left(\frac{\alpha+e}{e}x\right)^{1/\alpha} & \text{if } 0 \leq x < \frac{e}{\alpha+e}, \\ \ln\left[\frac{\alpha e}{(1-x)(\alpha+e)}\right] & \text{if } \frac{e}{\alpha+e} \leq x \leq 1, \end{cases}$$

for $\alpha \in (0,1)$.

3

Subsubproblem (b)

Using what we found in (a) we may use the inversion method to sample from $g(x)$ given in Equation (2), as shown in the code block beneath. Figure 2 shows one million samples drawn from `generate_from_gx()` and also the theoretical PDF.

```r
generate_from_gx <- function(n, alpha) {
  U <- runif(n)    # Generate n Uniform(0, 1) variables
  bound <- exp(1) / (alpha + exp(1))    # Boundary where G^(-1) changes
  left <- U < bound    # The left of the boundary
  U[left] <- (U[left] / bound)^(1 / alpha)    # Left CDF
  U[!left] <- 1 + log(alpha) - log(1 - U[!left]) - log(alpha + exp(1))    # Right CDF
  return(U)
}


# Sample
n <- 1000000  # One million samples
alpha <- 0.75
gx_samp <- generate_from_gx(n, alpha)    # Generating n samples from g(x)

# The theoretically correct PDF
theo_gx <- function(x, alpha) {
  const <- alpha * exp(1) / (alpha + exp(1))    # Normalizing constant
  func <- rep(0, length(x))    # Vector of zeros of same length as x
  left <- x > 0 & x < 1    # The PDF has one value for 0 < x < 1
  right <- x >= 1          # ... and one value for x >= 1
  func[left] <- const * x[left]^(alpha - 1)    # The value to the left
  func[right] <- const * exp(-x[right])        # The value to the right
  return(func)
}

# Plot
ggplot() +
  geom_histogram(
    data = as.data.frame(gx_samp),
    mapping = aes(x = gx_samp, y = ..density..),
    binwidth = 0.05,
    boundary = 0
  ) +
  stat_function(
    fun = theo_gx,
    args = list(alpha = alpha),
    aes(col = "Theoretical density")
  ) +
  xlim(0, 5) +
  ggtitle("Simulating from g(x) given in Equation (2)") +
  xlab("x") +
  ylab("Density") +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5))
```

Assuming $X \sim g(x)$ we may find the expectation to be

$$\mathrm{E}(X) = \int_{\mathbb{R}} x g(x)\, \mathrm{d}x = \int_0^1 c x^\alpha\, \mathrm{d}x + \int_1^\infty c x \mathrm{e}^{-x}\, \mathrm{d}x = \frac{\alpha \mathrm{e}}{(\alpha + 1)(\alpha + \mathrm{e})} + \frac{2\alpha}{\alpha + \mathrm{e}} \approx 0.768,$$
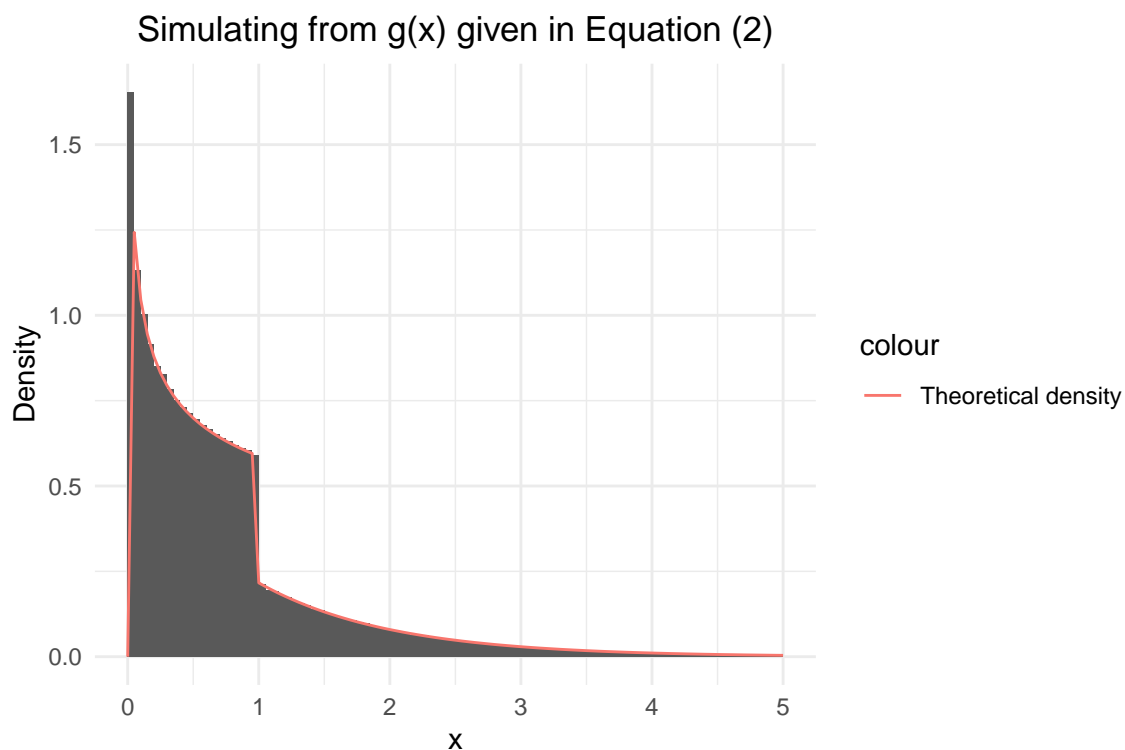
Figure 2: Normalized histogram of one million samples drawn from $g(x)$ given in Equation (2), together with the theoretical PDF, with $\alpha = 0.75$.

when $\alpha = 0.75$. This corresponds approximately to the sample mean shown in the following code block.

```
mean(gx_samp)
```

```
## [1] 0.76749
```

Similarly we may find the theoretical variance to be

$$\text{Var}(X) = \text{E}(X^2) - \text{E}(X)^2 \approx 0.705,$$

also for $\alpha = 0.75$, corresponding approximately to the sample variance given in the code block below.

```
var(gx_samp)
```

```
## [1] 0.703663
```

Subproblem 3.

Subsubproblem (a)

We consider the probability density function

$$f(x) = \frac{ce^{\alpha x}}{(1 + e^{\alpha x})^2},$$

for $-\infty < x < \infty$ and $\alpha > 0$. To find the normalizing constant c we make sure that the integral over $\mathbb{R}$ of $f(x)$ is one. That is

$$1 = \int_{\mathbb{R}} f(x)\,\mathrm{d}x = c \int_{\mathbb{R}} \frac{e^{\alpha x}}{(1 + e^{\alpha x})^2}\,\mathrm{d}x,$$

5

and letting $u = 1 + e^{\alpha x}$, it follows that

$$1 = \frac{c}{\alpha} \int_1^\infty \frac{\mathrm{d}u}{u^2} = \frac{c}{\alpha} \left[-\frac{1}{u} \right]_1^\infty = \frac{c}{\alpha}.$$

That is, the normalizing constant is $c = \alpha$, for $\alpha > 0$. We may then write the probability density function as

$$f(x) = \frac{\alpha e^{\alpha x}}{(1 + e^{\alpha x})^2}, \tag{3}$$

for $-\infty < x < \infty$ and $\alpha > 0$.

Subsubproblem (b)

The cumulative distribution function is given as

$$F(x) = \int_{-\infty}^x f(\xi)\,\mathrm{d}\xi = \int_{-\infty}^x \frac{\alpha e^{\alpha \xi}}{(1 + e^{\alpha \xi})^2}\,\mathrm{d}\xi.$$

Again letting $u = 1 + e^{\alpha \xi}$ it follows that

$$F(x) = \int_1^{1+e^{\alpha x}} \frac{\alpha e^{\alpha \xi}}{u^2} \frac{\mathrm{d}u}{\alpha e^{\alpha \xi}} = \int_1^{1+e^{\alpha x}} \frac{\mathrm{d}u}{u^2} = \left[\frac{1}{u} \right]_{1+e^{\alpha x}}^1 = 1 - \frac{1}{1 + e^{\alpha x}} = \frac{e^{\alpha x}}{1 + e^{\alpha x}},$$

which holds for $-\infty < x < \infty$ and $\alpha > 0$.

Solving $x = e^{\alpha y}/(1 + e^{\alpha y})$ for y then gives us the inverse cumulative distribution function. Some algebra then gives that

$$F^{-1}(x) = \frac{1}{\alpha} \ln\left(\frac{x}{1-x} \right) = \frac{1}{\alpha} \left[\ln(x) - \ln(1-x) \right],$$

for $0 < x < 1$ and $\alpha > 0$.

Subsubproblem (c)

Simulating from $f(x)$ given in Equation (3) is shown in the code block below, and the result is shown in Figure 3. The sampling is done by letting $U \sim \text{Uniform}(0, 1)$ and using inversion sampling.

```
generate_from_fx <- function(n, alpha) {
  U <- runif(n)    # Generate n Uniform(0, 1) variables
  X <- 1 / alpha * (log(U) - log(1 - U))    # Using the inverse CDF
  return(X)
}


# Sample
n <- 1000000  # One million samples
alpha <- 100  # Letting alpha be 100
fx_samp <- generate_from_fx(n, alpha)    # Generating n samples from f(x)

# The theoretically correct PDF
theo_fx <- function(x, alpha) {
  return(alpha * exp(alpha * x) / (1 + exp(alpha * x))^2)
}

# Plot
ggplot() +
  geom_histogram(
    data = as.data.frame(fx_samp),
```

```
  mapping = aes(x = fx_samp, y = ..density..),
  binwidth = 0.001,
  boundary = 0
) +
stat_function(
  fun = theo_fx,
  args = list(alpha = alpha),
  aes(col = "Theoretical density")
) +
ggtitle("Simulating from f(x) given in Equation (3)") +
xlab("x") +
ylab("Density") +
theme_minimal() +
theme(plot.title = element_text(hjust = 0.5))
```
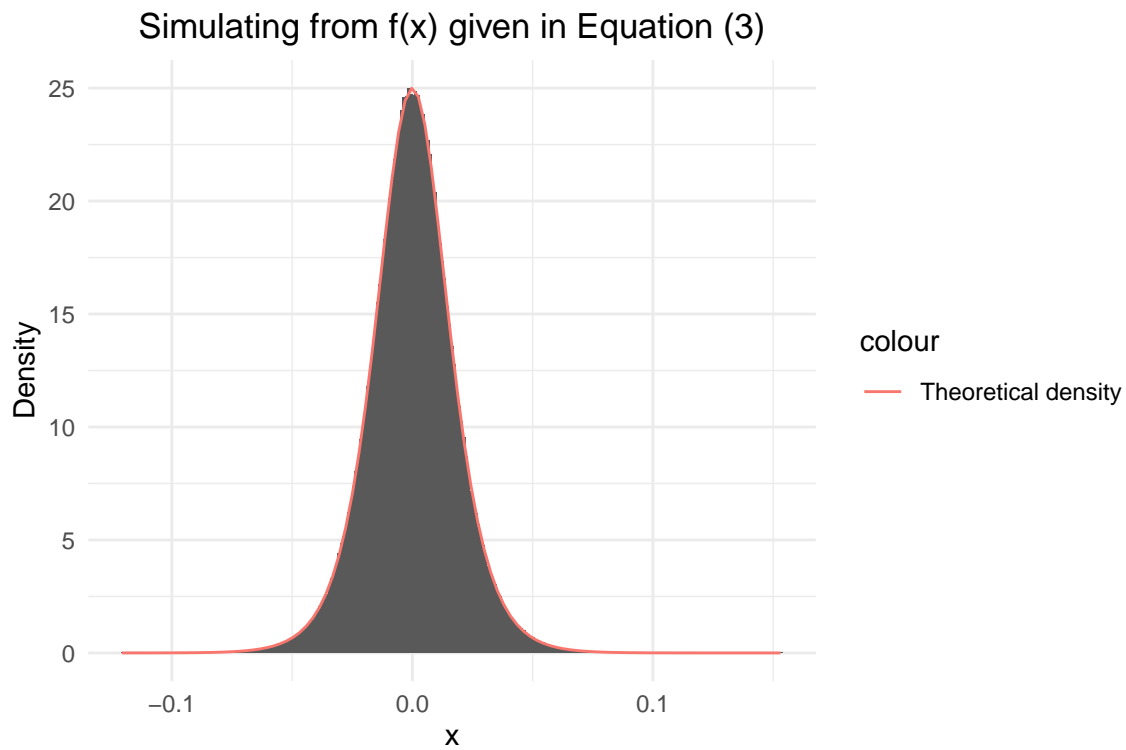

Figure 3: Normalized histogram of one million samples drawn from $f(x)$ given in Equation (3), together with the theoretical PDF, with $\alpha = 100$.

Letting $X \sim f(x)$ we may find the expected value to be

$$\mathrm{E}(X) = \int_{\mathbb{R}} x f(x)\, \mathrm{d}x = \int_{\mathbb{R}} \frac{\alpha x \mathrm{e}^{\alpha x}}{(1 + \mathrm{e}^{\alpha x})^2}\, \mathrm{d}x = 0,$$

because of symmetry. This is confirmed in the following code block where we see that the sample mean is approximately zero. This also holds for all $\alpha > 0$.

```
mean(fx_samp)
```

```
## [1] 2.957973e-05
```

We may also find the variance of X to be

$$\text{Var}(X) = \text{E}(X^2) = \int_{\mathbb{R}} \frac{\alpha x^2 e^{\alpha x}}{(1 + e^{\alpha x})^2} \, dx \approx 0.000329,$$

for $\alpha = 100$. This also corresponds to the sample variance as shown in the code block below.

```
var(fx_samp)
```

```
## [1] 0.0003288562
```

Subproblem 4.

We wish to simulate from a Normal$(0,1)$ distribution using the Box-Muller algoritm. If $X_1 \sim \text{Uniform}(0, \pi)$ and $X_2 \sim \text{Exponential}(1/2)$, then $Y_1 = \sqrt{X_2} \cos(X_1)$ and $Y_2 = \sqrt{X_2} \sin(X_1)$ are standard normal distributed. We use $Z = \sqrt{X_2} \cos(X_1)$ in the following code block. The result of the simulation can be seen in Figure 4, and it also shows the theoretical probability density function.

```r
std_normal <- function(n) {
  X1 <- pi * runif(n)    # n samples from Uniform(0, pi)
  X2 <- generate_from_exp(n, 1/2)    # n samples from Exponential(1/2)
  Z <- X2^(1/2) * cos(X1)    # Z ~ Normal(0, 1)
  return(Z)
}

# Sample
n <- 1000000   # One million samples
Box_Muller <- std_normal(n)    # Generating n samples from Normal(0, 1)

# Plot
ggplot() +
  geom_histogram(
    data = as.data.frame(Box_Muller),
    mapping = aes(x = Box_Muller, y = ..density..),
    binwidth = 0.05,
    boundary = 0
  ) +
  stat_function(
    fun = dnorm,
    aes(col = "Theoretical density")
  ) +
  ggtitle("Simulating from standard normal distribution") +
  xlab("z") +
  ylab("Density") +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5))
```

We know that if $Z \sim \text{Normal}(0,1)$, then $\text{E}(Z) = 0$ and $\text{Var}(Z) = 1$, and this corresponds to the approximate sample mean and variance shown in the code block below.

```
mean(Box_Muller)
```

```
## [1] -0.0006589413
```

```
var(Box_Muller)
```

```
## [1] 1.002119
```

Figure 4: One million samples drawn from a standard normal distribution using the Box-Muller algorithm, together with the theoretical PDF.

Subproblem 5.

We wish to simulate from a d-variate normal distribution with mean vector μ and covariance matrix Σ. Let $Z \sim \text{Normal}_d(0, I_d)$, where I_d is the identity matrix in $\mathbb{R}^{d \times d}$ and 0 is the zero-vector in $\mathbb{R}^d$. Then

$$X = \mu + AZ \sim \text{Normal}_d(\mu, AA^\top),$$

such that we need to find A such that $\Sigma = AA^\top$, this is done using `chol()` in R, and we construct the function in the following code block.

```
d_variate_normal <- function(n, mu, Sigma) {
  d <- length(mu)    # The dimension d is the dimension of mu
  A <- t(chol(Sigma)) # Cholesky decomposition of Sigma. Transpose to get lower triangular
  z <- std_normal(d * n)    # Create vector of d*n independent Normal(0, 1)
  Z <- matrix(z, nrow = d, ncol = n)    # Make z into (d X n) matrix Z
  X <- mu + A %*% Z    # X ~ Normal_d(mu, Sigma)
  return(X)
}
```

We now test the implementation using one million samples in $\mathbb{R}^3$, with

$$\mu = \begin{bmatrix} 1 \\ 7 \\ 2 \end{bmatrix} \quad \text{and} \quad \Sigma = \begin{bmatrix} 2 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 2 \end{bmatrix}.$$

We then expect the sample mean and sample covariance matrix to be approximately equal to these. In the following code block we see that this is indeed the case.

```r
# Sample
n <- 1000000      # One million samples
mu <- c(1, 7, 2)    # Create mu
Sigma <- matrix(c(2, -1, 0, -1, 2, -1, 0, -1, 2), nrow = 3)     # Create Sigma
normal <- d_variate_normal(n, mu, Sigma)     # Sample from d-variate normal

# Test
rowMeans(normal)  # Finding the mean of the rows of normal, sample mean
```

```
## [1] 1.000201 6.999770 1.997882
```

```r
cov(t(normal))     # Transpose because we want with respect to the rows
```

```
##                  [,1]        [,2]        [,3]
## [1,]   1.999879520 -0.9984456  0.003097756
## [2,]  -0.998445598  1.9978025 -0.998476727
## [3,]   0.003097756 -0.9984767  1.999870756
```

Problem B: The gamma distribution

1.

(a)

(b)

2.

(a)

(b)

3.

(a)

(b)

4.

5.

(a)

(b)

Problem C: Monte Carlo integration and variance reduction

1.

2.

3.

(a)

(b)

Problem D: Rejection sampling and importance sampling

1.

2.

3.

4.